

EDA Assignment

Group A

2/19/2025

Christian Reza

Goal:

- General schema description about an existing DB schema dataset (table/column count, relationships, etc.)
- Created & populated DB schema
- Learning points taken as we begin applying & fine-tuning T5 model to query DB schema (we do not expect to finish this, rather do technical exploration of a T5 model and how it can be fine-tuned)

Using spider data from Kaggle, specifically the soccer_1 data:

<https://www.kaggle.com/datasets/jeromeblanchet/yale-universitys-spider-10-nlp-dataset/data>

- I will be focusing on the **Player** table and create the training data off the following rows:
 - "player_name"
 - "birthday"
 - "height"
 - "weight"

Install packages and pull TAPAS QA model to fine-tune

```
In [1]: !pip install tf-keras
!pip install -q transformers
!pip install -q setuptools
!pip install -q tensorflow
!pip install -q pandas
import pandas as pd
from transformers import TapasConfig, TapasTokenizer, TFTapasForQuestionAnswering
import tensorflow as tf
```

```
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr
a8p0\LocalCache\local-packages\Python312\site-packages\tqdm\auto.py:21: TqdmWarning:
IPProgress not found. Please update jupyter and ipywidgets. See https://ipywidgets.re
adthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
```

WARNING:tensorflow:From C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\tf_keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.

```
In [2]: # The base sized model with WTQ configuration
model_name = "google/tapas-base-finetuned-wtq"
config = TapasConfig.from_pretrained(model_name)
model = TFTapasForQuestionAnswering.from_pretrained("google/tapas-base", config=config)
```

WARNING:tensorflow:From C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\tf_keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

All model checkpoint layers were used when initializing TFTapasForQuestionAnswering.

Some layers of TFTapasForQuestionAnswering were not initialized from the model checkpoint at google/tapas-base and are newly initialized: ['dropout_37', 'compute_column_logits', 'compute_token_logits', 'aggregation_classifier']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Some layers of TFTapasForQuestionAnswering were not initialized from the model checkpoint at google/tapas-base and are newly initialized: ['dropout_37', 'compute_column_logits', 'compute_token_logits', 'aggregation_classifier']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Create training data for TAPAS from Spider football data

Util script used supplied from *tapas_utils*:

https://github.com/NielsRogge/tapas_utils/blob/master/parse_answer_texts.py

```
In [3]: !pip install -q frozendict
!pip install -q scipy
!pip install -q numpy
```

```
In [4]: # coding=utf-8
# Copyright 2019 The Google AI Language Team Authors.
#
# Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
#     http://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
# Lint as: python3
"""This module implements a simple parser that can be used for TAPAS.
```

Given a table, a question and one or more answer_texts, it will parse the texts to populate other fields (e.g. answer_coordinates, float_value) that are required by TAPAS.

Please note that exceptions in this module are concise and not parameterized, since they are used as counter names in a BEAM pipeline.

"""

```
import enum
from typing import Callable, List, Text, Optional
```

```
import six
import struct
import unicodedata
import re
```

```
import frozendict
import numpy as np
import scipy.optimize
```

```
class SupervisionMode(enum.Enum):
    # Don't filter out any supervised information.
    NONE = 0
    # Remove all the supervised signals and recompute them by parsing answer
    # texts.
    REMOVE_ALL = 2
    # Same as above but discard ambiguous examples
    # (where an answer matches multiple cells).
    REMOVE_ALL_STRICT = 3
```

```
def _find_matching_coordinates(table, answer_text,
                               normalize):
    normalized_text = normalize(answer_text)
    for row_index, row in table.iterrows():
        for column_index, cell in enumerate(row):
            if normalized_text == normalize(str(cell)):
                yield (row_index, column_index)
```

```
def _compute_cost_matrix_inner(
    table,
    answer_texts,
    normalize,
    discard_ambiguous_examples,
):
```

"""Returns a cost matrix M where the value M[i,j] contains a matching cost from a

The matrix is a binary matrix and -1 is used to indicate a possible match from a given answer_texts to a specific cell table. The cost matrix can then be used to compute the optimal assignments that minimizes the cost using the hungarian algorithm (see scipy.optimize.linear_sum_assignment).

Args:

table: a Pandas dataframe.
answer_texts: a list of strings.
normalize: a function that normalizes a string.
discard_ambiguous_examples: If true discard if answer has multiple matches.

Raises:

ValueError if:

- we cannot correctly construct the cost matrix or the text-cell assignment is ambiguous.
- we cannot find a matching cell for a given answer_text.

Returns:

A numpy matrix with shape (num_answer_texts, num_rows * num_columns).

```
"""
max_candidates = 0
n_rows, n_columns = table.shape[0], table.shape[1]
num_cells = n_rows * n_columns
num_candidates = np.zeros((n_rows, n_columns))
cost_matrix = np.zeros((len(answer_texts), num_cells))

for index, answer_text in enumerate(answer_texts):
    found = 0
    for row, column in _find_matching_coordinates(table, answer_text,
                                                normalize):
        found += 1
        cost_matrix[index, (row * len(table.columns)) + column] = -1
        num_candidates[row, column] += 1
        max_candidates = max(max_candidates, num_candidates[row, column])
    if found == 0:
        return None
    if discard_ambiguous_examples and found > 1:
        raise ValueError("Found multiple cells for answers")

# TODO(piccinno): Shall we allow ambiguous assignments?
if max_candidates > 1:
    raise ValueError("Assignment is ambiguous")

return cost_matrix

def _compute_cost_matrix(
    table,
    answer_texts,
    discard_ambiguous_examples,
):
    """Computes cost matrix."""
    for index, normalize_fn in enumerate(STRING_NORMALIZATIONS):
        try:
            result = _compute_cost_matrix_inner(
                table,
                answer_texts,
                normalize_fn,
                discard_ambiguous_examples,
            )
            if result is None:
                continue
```

```

        return result
    except ValueError:
        if index == len(STRING_NORMALIZATIONS) - 1:
            raise
    return None

def _parse_answer_coordinates(table,
                              answer_texts,
                              discard_ambiguous_examples):
    """Populates answer_coordinates using answer_texts.

    Args:
        table: a Table message, needed to compute the answer coordinates.
        answer_texts: a list of strings
        discard_ambiguous_examples: If true discard if answer has multiple matches.

    Raises:
        ValueError if the conversion fails.
    """

    cost_matrix = _compute_cost_matrix(
        table,
        answer_texts,
        discard_ambiguous_examples,
    )
    if cost_matrix is None:
        return
    row_indices, column_indices = scipy.optimize.linear_sum_assignment(
        cost_matrix)

    # create answer coordinates as list of tuples
    answer_coordinates = []
    for row_index in row_indices:
        flatten_position = column_indices[row_index]
        row_coordinate = flatten_position // len(table.columns)
        column_coordinate = flatten_position % len(table.columns)
        answer_coordinates.append((row_coordinate, column_coordinate))

    return answer_coordinates

### START OF UTILITIES FROM TEXT_UTILS.PY ###

def wtq_normalize(x):
    """Returns the normalized version of x.
    This normalization function is taken from WikiTableQuestions github, hence the
    wtq prefix. For more information, see
    https://github.com/ppasupat/WikiTableQuestions/blob/master/evaluator.py
    Args:
        x: the object (integer type or string) to normalize.
    Returns:
        A normalized string.
    """
    x = x if isinstance(x, six.text_type) else six.text_type(x)
    # Remove diacritics.

```

```

x = "".join(
    c for c in unicodedata.normalize("NFKD", x)
    if unicodedata.category(c) != "Mn")
# Normalize quotes and dashes.
x = re.sub(u"['`^]", "'", x)
x = re.sub(u"['\""]", "'", x)
x = re.sub(u"[------]", "-", x)
x = re.sub(u"[-]", "", x)
while True:
    old_x = x
    # Remove citations.
    x = re.sub(u"(\?", "", x)
    x = x.replace("\n", " ")
    return x

_TOKENIZER = re.compile(r"\w+|[\^\w\s]+", re.UNICODE)

def tokenize_string(x):
    return list(_TOKENIZER.findall(x.lower()))

# List of string normalization functions to be applied in order. We go from
# simplest to more complex normalization procedures.
STRING_NORMALIZATIONS = (
    lambda x: x,
    lambda x: x.lower(),
    tokenize_string,
    wtq_normalize,
)

def to_float32(v):
    """If v is a float reduce precision to that of a 32 bit float."""
    if not isinstance(v, float):
        return v
    return struct.unpack("!f", struct.pack("!f", v))[0]

def convert_to_float(value):
    """Converts value to a float using a series of increasingly complex heuristics.
    Args:
    value: object that needs to be converted. Allowed types include
    float/int/strings.
    Returns:
    A float interpretation of value.
    Raises:
    ValueError if the float conversion of value fails.
    """
    if isinstance(value, float):
        return value
    if isinstance(value, int):
        return float(value)
    if not isinstance(value, six.string_types):
        raise ValueError("Argument value is not a string. Can't parse it as float")

```

```

sanitized = value

try:
    # Example: 1,000.7
    if "," in sanitized and "." in sanitized:
        return float(sanitized.replace(",", ""))
    # 1,000
    if "," in sanitized and _split_thousands(",", sanitized):
        return float(sanitized.replace(",", ""))
    # 5,5556
    if "," in sanitized and sanitized.count(",") == 1 and not _split_thousands(
        ",", sanitized):
        return float(sanitized.replace(",", "."))
    # 0.0.0.1
    if sanitized.count(".") > 1:
        return float(sanitized.replace(".", ""))
    # 0,0,0,1
    if sanitized.count(",") > 1:
        return float(sanitized.replace(",", ""))
    return float(sanitized)
except ValueError:
    # Avoid adding the sanitized value in the error message.
    raise ValueError("Unable to convert value to float")

### END OF UTILITIES FROM TEXT_UTILS.PY ###

def _parse_answer_float(answer_texts, float_value):
    if len(answer_texts) > 1:
        raise ValueError("Cannot convert to multiple answers to single float")
    float_value = convert_to_float(answer_texts[0])
    float_value = float_value

    return answer_texts, float_value

def _has_single_float_answer_equal_to(question, answer_texts, target):
    """Returns true if the question has a single answer whose value equals to target.
    if len(answer_texts) != 1:
        return False
    try:
        float_value = convert_to_float(answer_texts[0])
        # In general answer_float is derived by applying the same convert_to_float
        # function at interaction creation time, hence here we use exact match to
        # avoid any false positive.
        return to_float32(float_value) == to_float32(target)
    except ValueError:
        return False

def _parse_question(
    table,
    original_question,
    answer_texts,
    answer_coordinates,
    float_value,
    aggregation_function,

```

```

        clear_fields,
        discard_ambiguous_examples,
    ):
        """Parses question's answer_texts fields to possibly populate additional fields.

    Args:
        table: a Pandas dataframe, needed to compute the answer coordinates.
        original_question: a string.
        answer_texts: a list of strings, serving as the answer to the question.
        anser_coordinates:
        float_value: a float, serves as float value signal.
        aggregation_function:
        clear_fields: A list of strings indicating which fields need to be cleared
            and possibly repopulated.
        discard_ambiguous_examples: If true, discard ambiguous examples.

    Returns:
        A Question message with answer_coordinates or float_value field populated.

    Raises:
        ValueError if we cannot parse correctly the question message.
        """
    question = original_question

    # If we have a float value signal we just copy its string representation to
    # the answer text (if multiple answers texts are present OR the answer text
    # cannot be parsed to float OR the float value is different), after clearing
    # this field.
    if "float_value" in clear_fields and float_value is not None:
        if not _has_single_float_answer_equal_to(question, answer_texts, float_value):
            del answer_texts[:]
            float_value = float(float_value)
            if float_value.is_integer():
                number_str = str(int(float_value))
            else:
                number_str = str(float_value)
            answer_texts = []
            answer_texts.append(number_str)

    if not answer_texts:
        raise ValueError("No answer_texts provided")

    for field_name in clear_fields:
        if field_name == "answer_coordinates":
            answer_coordinates = None
        if field_name == "float_value":
            float_value = None
        if field_name == "aggregation_function":
            aggregation_function = None

    error_message = ""
    if not answer_coordinates:
        try:
            answer_coordinates = _parse_answer_coordinates(
                table,
                answer_texts,

```



```

        discard_ambiguous_examples,
    )
except ValueError as exc:
    error_message += "[answer_coordinates: {}]".format(str(exc))
    if discard_ambiguous_examples:
        raise ValueError(f"Cannot parse answer: {error_message}")

if not float_value:
    try:
        answer_texts, float_value = _parse_answer_float(answer_texts, float_value)
    except ValueError as exc:
        error_message += "[float_value: {}]".format(str(exc))

# Raises an exception if we cannot set any of the two fields.
if not answer_coordinates and not float_value:
    raise ValueError("Cannot parse answer: {}".format(error_message))

return question, answer_texts, answer_coordinates, float_value, aggregation_function

# TODO(piccinno): Use some sort of introspection here to get the field names of
# the proto.
_CLEAR_FIELDS = frozendict.frozendict({
    SupervisionMode.REMOVE_ALL: [
        "answer_coordinates", "float_value", "aggregation_function"
    ],
    SupervisionMode.REMOVE_ALL_STRICT: [
        "answer_coordinates", "float_value", "aggregation_function"
    ]
})

def parse_question(table, question, answer_texts, answer_coordinates=None, float_value=None,
                  mode=SupervisionMode.REMOVE_ALL):
    """Parses answer_text field of a question to populate additional fields required for supervision.

    Args:
        table: a Pandas dataframe, needed to compute the answer coordinates. Note that this function
        before supplying the table to this function.
        question: a string.
        answer_texts: a list of strings, containing one or more answer texts that should be used to
        answer_coordinates: optional answer coordinates supervision signal, if you have them.
        float_value: optional float supervision signal, if you already have this.
        aggregation_function: optional aggregation function supervised signal, if you have one.
        mode: see SupervisionMode enum for more information.

    Returns:
        A list with the question, populated answer_coordinates or float_value.

    Raises:
        ValueError if we cannot parse correctly the question string.
    """
    if mode == SupervisionMode.NONE:
        return question, answer_texts

    clear_fields = _CLEAR_FIELDS.get(mode, None)

```

```

if clear_fields is None:
    raise ValueError(f"Mode {mode.name} is not supported")

return _parse_question(
    table,
    question,
    answer_texts,
    answer_coordinates,
    float_value,
    aggregation_function,
    clear_fields,
    discard_ambiguous_examples=mode == SupervisionMode.REMOVE_ALL_STRICT,
)

```

```

In [70]: data = {
    "Player Name": [
        "Aaron Appindangoye",
        "Aaron Cresswell",
        "Aaron Doran",
        "Aaron Galindo",
        "Aaron Hughes",
        "Aaron Hunt",
        "Aaron Kuhl",
    ],
    "Birthday": [
        "1992-02-29 00:00:00",
        "1989-12-15 00:00:00",
        "1991-05-13 00:00:00",
        "1982-05-08 00:00:00",
        "1979-11-08 00:00:00",
        "1986-09-04 00:00:00",
        "1996-01-30 00:00:00",
    ],
    "Weight": [
        "187",
        "146",
        "163",
        "198",
        "154",
        "161",
        "146",
    ],
    "Height": [
        "182.88",
        "170.18",
        "170.18",
        "182.88",
        "182.88",
        "182.88",
        "172.72",
    ],
}

```

```

table = pd.DataFrame.from_dict(data)
table = table.astype(str)

```

```
table.head(8)
```

Out[70]:

	Player Name	Birthday	Weight	Height
0	Aaron Appindangoye	1992-02-29 00:00:00	187	182.88
1	Aaron Cresswell	1989-12-15 00:00:00	146	170.18
2	Aaron Doran	1991-05-13 00:00:00	163	170.18
3	Aaron Galindo	1982-05-08 00:00:00	198	182.88
4	Aaron Hughes	1979-11-08 00:00:00	154	182.88
5	Aaron Hunt	1986-09-04 00:00:00	161	182.88
6	Aaron Kuhl	1996-01-30 00:00:00	146	172.72

In [71]:

```
question = "How much does Aaron Hughes weigh?"

answer_texts = ["154"]

question, answer_texts, answer_coordinates, float_value, aggregation_function = par
print(question)
print(answer_texts)
print("Found coordinates:", answer_coordinates)
print("Found float value:", float_value)
```

How much does Aaron Hughes weigh?

['154']

Found coordinates: [(np.int64(4), np.int64(2))]

Found float value: 154.0

In [72]:

```
question = "What is Aaron Appindangoye's birthday?"

answer_texts = ["1992-02-29 00:00:00"]

question, answer_texts, answer_coordinates, float_value, aggregation_function = par
print(question)
print(answer_texts)
print("Found coordinates:", answer_coordinates)
print("Found float value:", float_value)
```

What is Aaron Appindangoye's birthday?

['1992-02-29 00:00:00']

Found coordinates: [(np.int64(0), np.int64(1))]

Found float value: None

In [73]:

```
question = "How many players weigh more than 150?"

answer_texts = ["187", "163", "198", "154", "161"]
float_value = 4

question, answer_texts, answer_coordinates, float_value, aggregation_function = par
print(question)
print(answer_texts)
```

```
print("Found coordinates:", answer_coordinates)
print("Found float value:", float_value)
```

How many players weigh more than 150?

['187', '163', '198', '154', '161']

Found coordinates: [(np.int64(0), np.int64(2)), (np.int64(2), np.int64(2)), (np.int64(3), np.int64(2)), (np.int64(4), np.int64(2)), (np.int64(5), np.int64(2))]

Found float value: None

Take output from SQA Util script to create tensor data to fine-tune model

In [9]: `tokenizer = TapasTokenizer.from_pretrained(model_name)`

```
data = {
    "Player Name": [
        "Aaron Appindangoye",
        "Aaron Cresswell",
        "Aaron Doran",
        "Aaron Galindo",
        "Aaron Hughes",
        "Aaron Hunt",
        "Aaron Kuhl",
    ],
    "Birthday": [
        "1992-02-29 00:00:00",
        "1989-12-15 00:00:00",
        "1991-05-13 00:00:00",
        "1982-05-08 00:00:00",
        "1979-11-08 00:00:00",
        "1986-09-04 00:00:00",
        "1996-01-30 00:00:00",
    ],
    "Weight": [
        "187",
        "146",
        "163",
        "198",
        "154",
        "161",
        "146",
    ],
    "Height": [
        "182.88",
        "170.18",
        "170.18",
        "182.88",
        "182.88",
        "182.88",
        "172.72",
    ],
}
queries = [
    "How much does Aaron Hughes weigh?",
    "What is Aaron Appindangoye's birthday?",
```

```

    "How many players weigh more than 150?",
]
answer_coordinates = [[(4, 2)], [(0, 1)], [(0, 2), (2, 2), (3, 2), (4, 2), (5, 2)]]
answer_text = [["154"], ["1992-02-29 00:00:00"], ["4"]]
table = pd.DataFrame.from_dict(data)
inputs = tokenizer(
    table=table,
    queries=queries,
    answer_coordinates=answer_coordinates,
    answer_text=answer_text,
    padding="max_length",
    return_tensors="tf",
)
inputs

```

C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok
enization_tapas.py:2699: FutureWarning: Series.__getitem__ treating keys as position
s is deprecated. In a future version, integer keys will always be treated as labels
(consistent with DataFrame behavior). To access a value by position, use `ser.iloc[p
os]`

```
text = normalize_for_match(row[col_index].text)
```

C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok
enization_tapas.py:1493: FutureWarning: Series.__getitem__ treating keys as position
s is deprecated. In a future version, integer keys will always be treated as labels
(consistent with DataFrame behavior). To access a value by position, use `ser.iloc[p
os]`

```
cell = row[col_index]
```

```

Out[9]: {'input_ids': <tf.Tensor: shape=(3, 512), dtype=int32, numpy=
array([[ 101, 2129, 2172, ..., 0, 0, 0],
       [ 101, 2054, 2003, ..., 0, 0, 0],
       [ 101, 2129, 2116, ..., 0, 0, 0]], dtype=int32)>, 'labels': <tf.Tensor: shape=(3, 512), dtype=int32, numpy=
array([[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]], dtype=int32)>, 'numeric_values': <tf.Tensor: shape=(3, 512), dtype=float32, numpy=
array([[nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan],
       [nan, nan, nan, ..., nan, nan, nan]], dtype=float32)>, 'numeric_values_scale': <tf.Tensor: shape=(3, 512), dtype=float32, numpy=
array([[1., 1., 1., ..., 1., 1., 1.],
       [1., 1., 1., ..., 1., 1., 1.],
       [1., 1., 1., ..., 1., 1., 1.]], dtype=float32)>, 'token_type_ids': <tf.Tensor: shape=(3, 512, 7), dtype=int32, numpy=
array([[[[0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         ...,
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0]],
        ...,
        [[0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         ...,
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0]],
        ...,
        [[0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         ...,
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0],
         [0, 0, 0, ..., 0, 0, 0]]], dtype=int32)>, 'attention_mask': <tf.Tensor: shape=(3, 512), dtype=int32, numpy=
array([[1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0]], dtype=int32)>}]

```

```

In [67]: import os
from pathlib import Path

sqa_path = Path("C:/School/Spring 2025/DSCI-2025-TEAM-A/reza/data/sqa_train_set.csv")
table_csv_path = Path("C:/School/Spring 2025/DSCI-2025-TEAM-A/reza/data/")

class TableDataset:
    """
    A dataset class for handling table data and tokenizing it for model input.
    Attributes:

```

data (pd.DataFrame): The input data containing table file paths, questions,
 tokenizer (Tokenizer): The tokenizer used to process the table data and que
 Methods:
 __iter__():
 Iterates over the dataset, tokenizes the table data and questions, and
 __len__():
 Returns the length of the dataset.
 """

```
def __init__(self, data, tokenizer):
    self.data = data
    self.tokenizer = tokenizer

def __iter__(self):
    for idx in range(self.__len__()):
        item = self.data.iloc[idx]
        table = pd.read_csv(str(table_csv_path) + os.sep + item.table_file).ast
            str
        ) # be sure to make your table data text only
        encoding = self.tokenizer(
            table=table,
            queries=item.question,
            answer_coordinates=eval(item.answer_coordinates),
            # answer_coordinates=item.answer_coordinates,
            answer_text=item.answer_text,
            truncation=True,
            padding="max_length",
            return_tensors="tf",
        )
        # remove the batch dimension which the tokenizer adds by default
        encoding = {key: tf.squeeze(val, 0) for key, val in encoding.items()}
        # add the float_answer which is also required (weak supervision for agg
        float_answer = item.float_answer if not pd.isna(item.float_answer) else
        encoding["float_answer"] = tf.convert_to_tensor(
            item.float_answer, dtype=tf.float32
        )
        yield encoding["input_ids"], encoding["attention_mask"], encoding[
            "numeric_values"
        ], encoding["numeric_values_scale"], encoding["token_type_ids"], encodi
            "labels"
        ], encoding[
            "float_answer"
        ]

def __len__(self):
    return len(self.data)
```

```
data = pd.read_csv(sqa_path, sep="\t")
train_dataset = TableDataset(data, tokenizer)
output_signature = (
    tf.TensorSpec(shape=(512,), dtype=tf.int32),
    tf.TensorSpec(shape=(512,), dtype=tf.int32),
    tf.TensorSpec(shape=(512,), dtype=tf.float32),
    tf.TensorSpec(shape=(512,), dtype=tf.float32),
    tf.TensorSpec(shape=(512, 7), dtype=tf.int32),
```

```

tf.TensorSpec(shape=(512,), dtype=tf.int32),
tf.TensorSpec(shape=(), dtype=tf.float32),
)
train_dataloader = tf.data.Dataset.from_generator(
    lambda: iter(train_dataset), output_signature=output_signature
).batch(32)

```

Fine-tuning the TAPAS model with new data!

- This I am confused on, what was the purpose of creating the `inputs` from the dataframe, do we still need the tsv to train the data?
- What is wrong with my `train_dataloader` that is causing an error in the batches - is there something I missed in the data loader, do I need more training data?

```

In [68]: # this is the default WTQ configuration
config = TapasConfig(
    num_aggregation_labels=4,
    use_answer_as_supervision=True,
    answer_loss_cutoff=0.664694,
    cell_selection_preference=0.207951,
    huber_loss_delta=0.121194,
    init_cell_selection_weights_to_zero=True,
    select_one_column=True,
    allow_empty_column_selection=False,
    temperature=0.0352513,
)
model = TFTapasForQuestionAnswering.from_pretrained(model_name, config=config)

optimizer = tf.keras.optimizers.Adam(learning_rate=5e-5)

for epoch in range(2): # Loop over the dataset multiple times
    for batch in train_dataloader:
        # get the inputs;
        input_ids = batch[0]
        attention_mask = batch[1]
        token_type_ids = batch[4]
        labels = batch[-1]
        numeric_values = batch[2]
        numeric_values_scale = batch[3]
        float_answer = batch[6]

        # forward + backward + optimize
        with tf.GradientTape() as tape:
            outputs = model(
                input_ids=input_ids,
                attention_mask=attention_mask,
                token_type_ids=token_type_ids,
                labels=labels,
                numeric_values=numeric_values,
                numeric_values_scale=numeric_values_scale,
                float_answer=float_answer,
            )

```



```
grads = tape.gradient(outputs.loss, model.trainable_weights)
optimizer.apply_gradients(zip(grads, model.trainable_weights))
```

```
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok  
enization_tapas.py:2699: FutureWarning: Series.__getitem__ treating keys as position  
s is deprecated. In a future version, integer keys will always be treated as labels  
(consistent with DataFrame behavior). To access a value by position, use `ser.iloc[p  
os]`
```

```
text = normalize_for_match(row[col_index].text)
```

```
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok  
enization_tapas.py:1493: FutureWarning: Series.__getitem__ treating keys as position  
s is deprecated. In a future version, integer keys will always be treated as labels  
(consistent with DataFrame behavior). To access a value by position, use `ser.iloc[p  
os]`
```

```
cell = row[col_index]
```

Some layers from the model checkpoint at google/tapas-base-finetuned-wtq were not used when initializing TFTapasForQuestionAnswering: ['dropout_570']

- This IS expected if you are initializing TFTapasForQuestionAnswering from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

- This IS NOT expected if you are initializing TFTapasForQuestionAnswering from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Some layers of TFTapasForQuestionAnswering were not initialized from the model checkpoint at google/tapas-base-finetuned-wtq and are newly initialized: ['dropout_417']

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok  
enization_tapas.py:2699: FutureWarning: Series.__getitem__ treating keys as position  
s is deprecated. In a future version, integer keys will always be treated as labels  
(consistent with DataFrame behavior). To access a value by position, use `ser.iloc[p  
os]`
```

```
text = normalize_for_match(row[col_index].text)
```

```
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok  
enization_tapas.py:1493: FutureWarning: Series.__getitem__ treating keys as position  
s is deprecated. In a future version, integer keys will always be treated as labels  
(consistent with DataFrame behavior). To access a value by position, use `ser.iloc[p  
os]`
```

```
cell = row[col_index]
```

```
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok  
enization_tapas.py:2699: FutureWarning: Series.__getitem__ treating keys as position  
s is deprecated. In a future version, integer keys will always be treated as labels  
(consistent with DataFrame behavior). To access a value by position, use `ser.iloc[p  
os]`
```

```
text = normalize_for_match(row[col_index].text)
```

```
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok  
enization_tapas.py:1493: FutureWarning: Series.__getitem__ treating keys as position  
s is deprecated. In a future version, integer keys will always be treated as labels  
(consistent with DataFrame behavior). To access a value by position, use `ser.iloc[p  
os]`
```

```
cell = row[col_index]
```

```
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr  
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok
```

```
enization_tapas.py:2699: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
    text = normalize_for_match(row[col_index].text)
C:\Users\rza_t\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfr
a8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\tok
enization_tapas.py:1493: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future version, integer keys will always be treated as labels (consistent with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`
    cell = row[col_index]
```

InvalidArgumentError

Traceback (most recent call last)

Cell In[68], line 30

```
28 # forward + backward + optimize
29 with tf.GradientTape() as tape:
--> 30     outputs = model(
31         input_ids=input_ids,
32         attention_mask=attention_mask,
33         token_type_ids=token_type_ids,
34         labels=labels,
35         numeric_values=numeric_values,
36         numeric_values_scale=numeric_values_scale,
37         float_answer=float_answer,
38     )
39 grads = tape.gradient(outputs.loss, model.trainable_weights)
40 optimizer.apply_gradients(zip(grads, model.trainable_weights))
```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\tf_keras\src\utils\traceback_utils.py:70, in filter_traceback.<locals>.error_handler(*args, **kwargs)

```
67     filtered_tb = _process_traceback_frames(e.__traceback__)
68     # To get the full stack trace, call:
69     # `tf.debugging.disable_traceback_filtering()`
--> 70     raise e.with_traceback(filtered_tb) from None
71 finally:
72     del filtered_tb
```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\transformers\modeling_tf_utils.py:437, in unpack_inputs.<locals>.run_call_with_unpacked_inputs(self, *args, **kwargs)

```
434     config = self.config
436     unpacked_inputs = input_processing(func, config, **fn_args_and_kwargs)
--> 437     return func(self, **unpacked_inputs)
```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\modeling_tf_tapas.py:1465, in TFTapasForQuestionAnswering.call(self, input_ids, attention_mask, token_type_ids, position_ids, head_mask, inputs_embeds, table_mask, aggregation_labels, float_answer, numeric_values, numeric_values_scale, output_attentions, output_hidden_states, return_dict, labels, training)

```
1391 @unpack_inputs
1392 @add_start_docstrings_to_model_forward(TAPAS_INPUTS_DOCSTRING.format("batch_size, sequence_length"))
1393 @replace_return_docstrings(output_type=TFTableQuestionAnsweringOutput, config_class=_CONFIG_FOR_DOC)
1394 (...)
1411     training: Optional[bool] = False,
1412 ) -> Union[TFTableQuestionAnsweringOutput, Tuple[tf.Tensor]]:
1413     """
1414     table_mask (`tf.Tensor` of shape `(batch_size, seq_length)`, *optional*):
1415         Mask for the table. Indicates which tokens belong to the table (1).
1416         Question tokens, table headers and
1417         (...)
1462     >>> logits_aggregation = outputs.logits_aggregation
1463     """
```

```

-> 1465     outputs = self.tapas(
1466         input_ids=input_ids,
1467         attention_mask=attention_mask,
1468         token_type_ids=token_type_ids,
1469         position_ids=position_ids,
1470         head_mask=head_mask,
1471         inputs_embeds=inputs_embeds,
1472         output_attentions=output_attentions,
1473         output_hidden_states=output_hidden_states,
1474         return_dict=return_dict,
1475         training=training,
1476     )
1478     sequence_output = outputs[0]
1479     pooled_output = outputs[1]

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\transformers\modeling_tf_utils.py:437, in `unpack_inputs.<locals>.run_call_with_unpacked_inputs(self, *args, **kwargs)`

```

434     config = self.config
436     unpacked_inputs = input_processing(func, config, **fn_args_and_kwargs)
-> 437     return func(self, **unpacked_inputs)

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\modeling_tf_tapas.py:888, in `TFTapasMainLayer.call(self, input_ids, attention_mask, token_type_ids, position_ids, head_mask, inputs_embeds, output_attentions, output_hidden_states, return_dict, training)`

```

885     if token_type_ids is None:
886         token_type_ids = tf.fill(dims=input_shape + [len(self.config.type_vocab_sizes)], value=0)
-> 888     embedding_output = self.embeddings(
889         input_ids=input_ids,
890         position_ids=position_ids,
891         token_type_ids=token_type_ids,
892         inputs_embeds=inputs_embeds,
893         training=training,
894     )
896     # We create a 3D attention mask from a 2D tensor mask.
897     # Sizes are [batch_size, 1, 1, to_seq_length]
898     # So we can broadcast to [batch_size, num_heads, from_seq_length, to_seq_length]
899     # this attention mask is more simple than the triangular masking of causal attention
900     # used in OpenAI GPT, we just need to prepare the broadcast dimension here.
901     extended_attention_mask = tf.reshape(attention_mask, (input_shape[0], 1, 1, input_shape[1]))

```

File ~\AppData\Local\Packages\PythonSoftwareFoundation.Python.3.12_qbz5n2kfra8p0\LocalCache\local-packages\Python312\site-packages\transformers\models\tapas\modeling_tf_tapas.py:226, in `TFTapasEmbeddings.call(self, input_ids, position_ids, token_type_ids, inputs_embeds, training)`

```

224     for i in range(self.number_of_token_type_embeddings):
225         name = f"token_type_embeddings_{i}"
-> 226         final_embeddings += tf.gather(params=getattr(self, name), indices=token_type_ids[:, :, i])
228     final_embeddings = self.LayerNorm(inputs=final_embeddings)

```

```
229 final_embeddings = self.dropout(inputs=final_embeddings, training=training)
```

InvalidArgumentError: Exception encountered when calling layer 'embeddings' (type TFTapasEmbeddings).

{{function_node __wrapped__ResourceGather_device_/job:localhost/replica:0/task:0/device:CPU:0}} indices[1,160] = 5341 is not in [0, 256) [Op:ResourceGather] name:

Call arguments received by layer 'embeddings' (type TFTapasEmbeddings):

- input_ids=tf.Tensor(shape=(3, 512), dtype=int32)
- position_ids=None
- token_type_ids=tf.Tensor(shape=(3, 512, 7), dtype=int32)
- inputs_embeds=None
- training=False

Use the newly fine-tune model against soccer dataset!

TBD